

Lecture 8: Nonlinearity (Polynomial Regression, Spline & GAM)

BIOS635

02/04/2020

Moving beyond linearity

The truth is never linear!

Or almost never!

But often the linearity assumption is good enough.

When its not ...

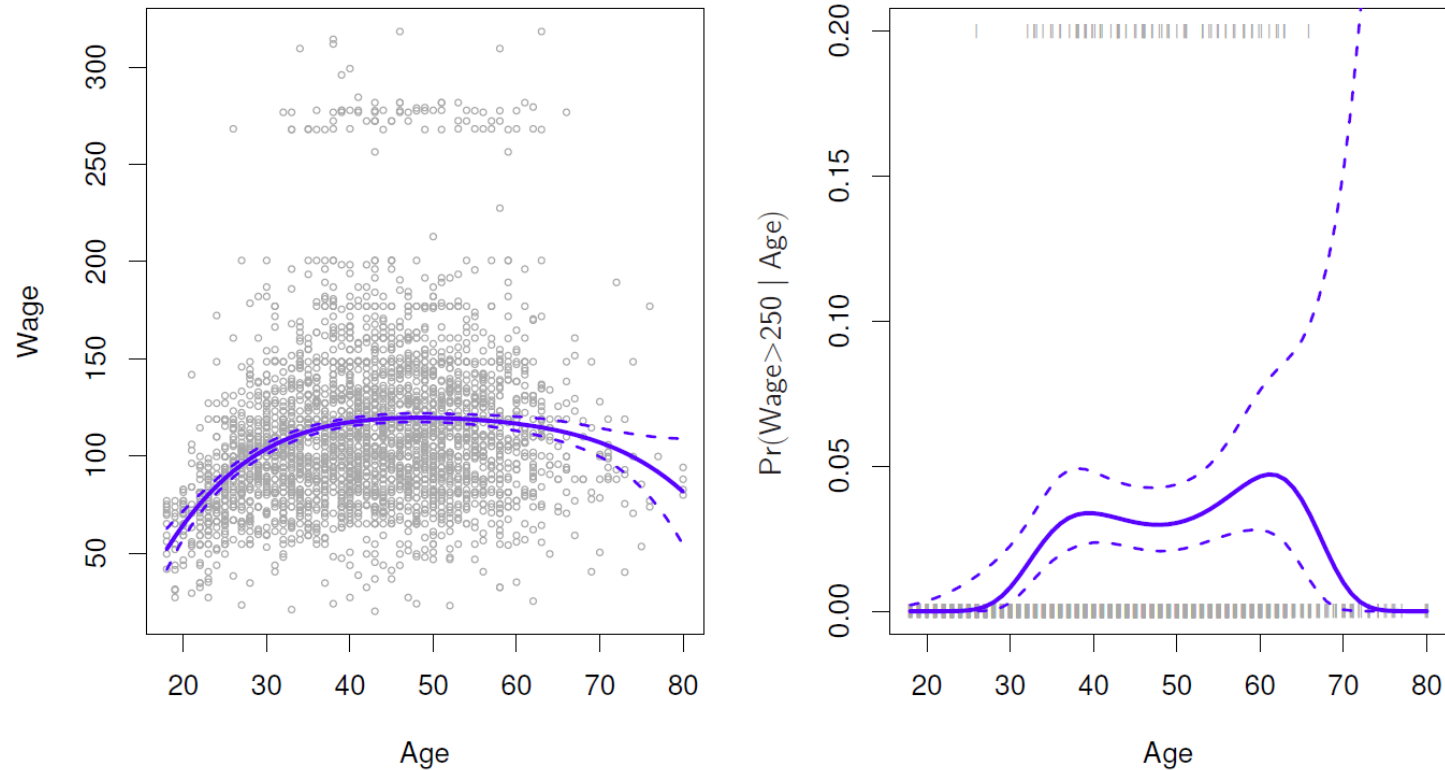
- polynomials,
- step functions,
- splines,
- local regression, and
- generalized additive models

offer a lot of flexibility, without losing the ease and interpretability of linear models.

Polynomial regression

$$y_i = \beta_0 + \beta_1 x_i + \beta_2 x_i^2 + \beta_3 x_i^3 + \dots + \beta_d x_i^d + \epsilon_i$$

Degree-4 Polynomial



Polynomial regression: linear

- Create new variables $X_1 = X$, $X_2 = X^2$, etc and then treat as multiple linear regression.
- Not really interested in the coefficients; more interested in the fitted function values at any value x_0 :

$$\hat{f}(x_0) = \hat{\beta}_0 + \hat{\beta}_1 x_0 + \hat{\beta}_2 x_0^2 + \hat{\beta}_3 x_0^3 + \hat{\beta}_4 x_0^4.$$


- Since $\hat{f}(x_0)$ is a linear function of the $\hat{\beta}_\ell$, can get a simple expression for *pointwise-variances* $\text{Var}[\hat{f}(x_0)]$ at any value x_0 . In the figure we have computed the fit and pointwise standard errors on a grid of values for x_0 . We show $\hat{f}(x_0) \pm 2 \cdot \text{se}[\hat{f}(x_0)]$.
- We either fix the degree d at some reasonably low value, else use cross-validation to choose d .

Polynomial regression: logistic


- Logistic regression follows naturally. For example, in figure we model

$$\Pr(y_i > 250|x_i) = \frac{\exp(\beta_0 + \beta_1 x_i + \beta_2 x_i^2 + \dots + \beta_d x_i^d)}{1 + \exp(\beta_0 + \beta_1 x_i + \beta_2 x_i^2 + \dots + \beta_d x_i^d)}.$$


- To get confidence intervals, compute upper and lower bounds on *on the logit scale*, and then invert to get on probability scale.
- Can do separately on several variables—just stack the variables into one matrix, and separate out the pieces afterwards (see GAMs later).
- Caveat: polynomials have notorious tail behavior — very bad for extrapolation.
- Can fit using $\mathbf{y} \sim \text{poly}(\mathbf{x}, \text{degree} = 3)$ in formula.




`lm(wage ~ age, data = Wage)`



`lm(wage ~ poly(age, 2), data = Wage)`



`lm(wage ~ poly(age, 4), data = Wage)`



`lm(wage ~ poly(age, 10), data = Wage)`

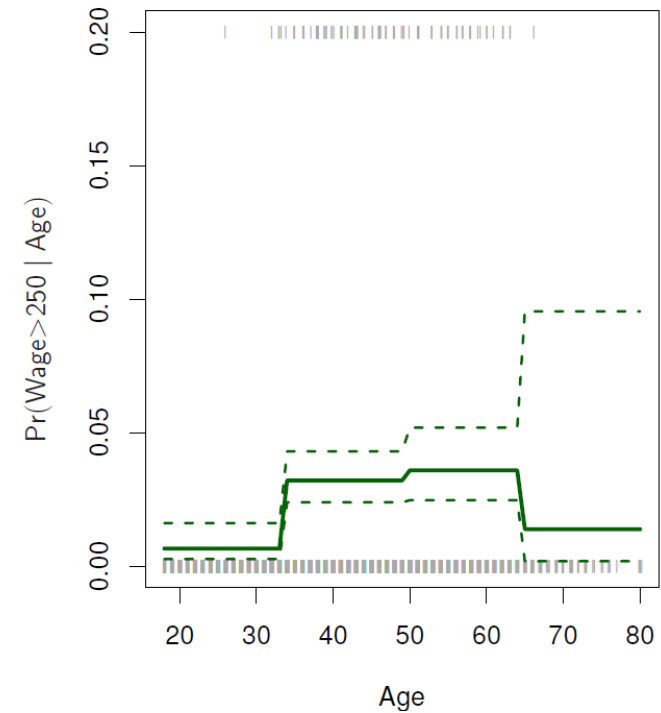

Step functions

Another way of creating transformations of a variable — cut the variable into distinct regions.

$$C_1(X) = I(X < 35), \quad C_2(X) = I(35 \leq X < 65), \dots, C_3(X) = I(X \geq 65)$$

Piecewise Constant

Age	C_1	C_2	C_3
18	1	0	0
24	1	0	0
45	0	1	0
67	0	0	1
54	0	1	0
⋮	⋮	⋮	⋮



Step functions, cont'd

- Easy to work with. Creates a series of dummy variables representing each group.
- Useful way of creating interactions that are easy to interpret. For example, interaction effect of **Year** and **Age**:

$$I(\text{Year} < 2005) \cdot \text{Age}, \quad I(\text{Year} \geq 2005) \cdot \text{Age}$$

would allow for different linear functions in each age category.


- In R: `I(year < 2005)` or `cut(age, c(18, 25, 40, 65, 90))`.
- Choice of cutpoints or *knots* can be problematic. For creating nonlinearities, smoother alternatives such as *splines* are available.

Piecewise polynomials


- Instead of a single polynomial in X over its whole domain, we can rather use different polynomials in regions defined by knots. E.g. (see figure)

$$y_i = \begin{cases} \beta_{01} + \beta_{11}x_i + \beta_{21}x_i^2 + \beta_{31}x_i^3 + \epsilon_i & \text{if } x_i < c; \\ \beta_{02} + \beta_{12}x_i + \beta_{22}x_i^2 + \beta_{32}x_i^3 + \epsilon_i & & \text{if } x_i \geq c. \end{cases}$$


- Better to add constraints to the polynomials, e.g. continuity.
- *Splines* have the “maximum” amount of continuity.




Piecewise Cubic




Continuous Piecewise Cubic



Cubic Spline



Linear Spline



Linear splines

A linear spline with knots at ξ_k , $k = 1, \dots, K$ is a piecewise linear polynomial continuous at each knot.

We can represent this model as



$$y_i = \beta_0 + \beta_1 b_1(x_i) + \beta_2 b_2(x_i) + \dots + \beta_{K+3} b_{K+3}(x_i) + \epsilon_i,$$

where the b_k are *basis functions*.

$$\begin{aligned} b_1(x_i) &= x_i \\ b_{k+1}(x_i) &= (x_i - \xi_k)_+, \quad k = 1, \dots, K \end{aligned}$$

Here the $()_+$ means *positive part*; i.e.

$$(x_i - \xi_k)_+ = \begin{cases} x_i - \xi_k & \text{if } x_i > \xi_k \\ 0 & \text{otherwise} \end{cases}$$



Cubic splines

A cubic spline with knots at ξ_k , $k = 1, \dots, K$ is a piecewise cubic polynomial with continuous derivatives up to order 2 at each knot.

Again we can represent this model with truncated power basis functions

$$y_i = \beta_0 + \beta_1 b_1(x_i) + \beta_2 b_2(x_i) + \cdots + \beta_{K+3} b_{K+3}(x_i) + \epsilon_i,$$

$$b_1(x_i) = x_i$$



$$b_2(x_i) = x_i^2$$

$$b_3(x_i) = x_i^3$$

$$b_{k+3}(x_i) = (x_i - \xi_k)_+^3, \quad k = 1, \dots, K$$

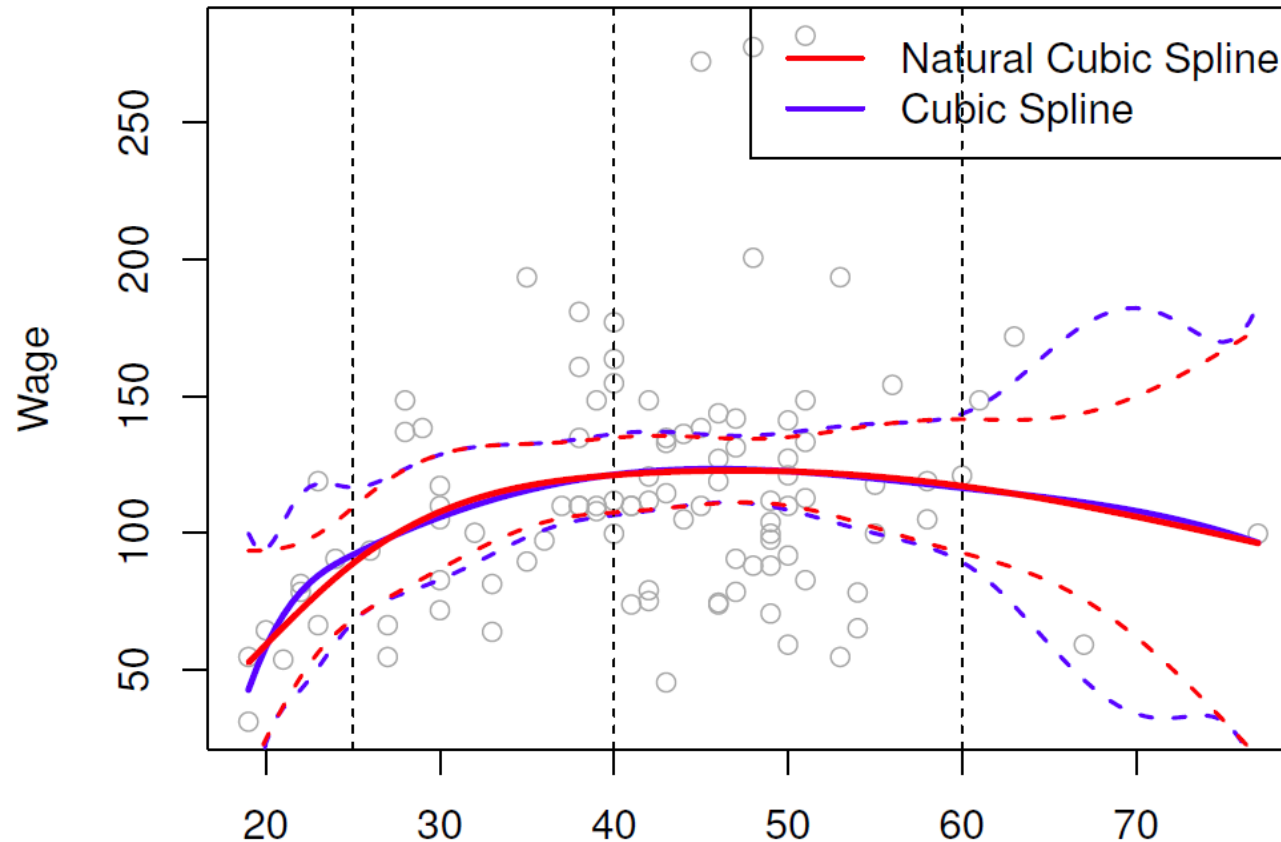
where

$$(x_i - \xi_k)_+^3 = \begin{cases} (x_i - \xi_k)^3 & \text{if } x_i > \xi_k \\ 0 & \text{otherwise} \end{cases}$$



Natural cubic splines


Second derivatives of the spline polynomials are set equal to zero at the end points of interpolation. This forces the spline to be a straight line outside the interval.



Fitting splines in R


Fitting splines in R is easy: `bs(x, ...)` for any degree splines, and `ns(x, ...)` for natural cubic splines, in package `splines`.

Natural Cubic Spline



Knot placement & degree of freedom

- One strategy is to decide K , the number of knots, and then place them at appropriate quantiles of the observed X .
- A cubic spline with K knots has $K + 4$ parameters or degrees of freedom.
- A natural spline with K knots has K degrees of freedom.



Comparison of a degree-14 polynomial and a natural cubic spline, each with 15df.

```
ns(age, df=14)
```

```
poly(age, deg=14)
```

Smoothing splines

- The **smoothing spline** estimator is the solution \hat{g} to the problem

$$\text{minimize } \underbrace{\sum_{i=1}^n (y_i - g(x_i))^2}_{\text{RSS}} + \underbrace{\lambda \int g''(t)^2 dt}_{\text{Roughness penalty}}$$

- This is a **penalized regression problem**
- We're saying we want a function that:
 - ① Fits the data well; and
 - ② isn't too *wiggly*
- Large $\lambda \implies \hat{g}$ will have low variability (& higher bias)
- Small $\lambda \implies \hat{g}$ will have high variability (& lower bias)

How is this *at all* related to splines?

Smoothing splines, cont'd


$$\text{minimize } \underbrace{\sum_{i=1}^n (y_i - g(x_i))^2}_{\text{RSS}} + \underbrace{\lambda \int g''(t)^2 dt}_{\text{Roughness penalty}} \quad (*)$$

It turns out...


- The solution to $(*)$ is a **natural cubic spline**
- The solution has knots at every unique value of x
- The **effective degrees of freedom** of the solution is calculable
- $\lambda \longleftrightarrow df$

Specify df in R: `smooth.spline(age, wage, df = 10)`

Choose lambda via leave-one-out cross validation (LOOCV) in R: `smooth.spline(age, wage)`



Local regression



With a sliding weight function, we fit separate linear fits over the range of X by weighted least squares.

See text for more details, and `loess()` function in R.

Nonlinearity coding in R

Model	R command
Degree-d polynomial regression	<code>~ poly(x, degree = d)</code>
Step functions with knots c1, c2, c3	<code>~ cut(x, breaks = c(c1, c2, c3))</code>
Cubic spline	<code>~ bs(x, df)</code>
Natural cubic spline	<code>~ ns(x, df)</code>
Degree-d spline	<code>~ bs(x, df, degree = d)</code>
Smoothing spline	<code>~ s(x, df)</code>
Local linear regression	<code>~ lo(x)</code>

Generalized additive models (GAMs)

- Recall the **Linear Regression Model**

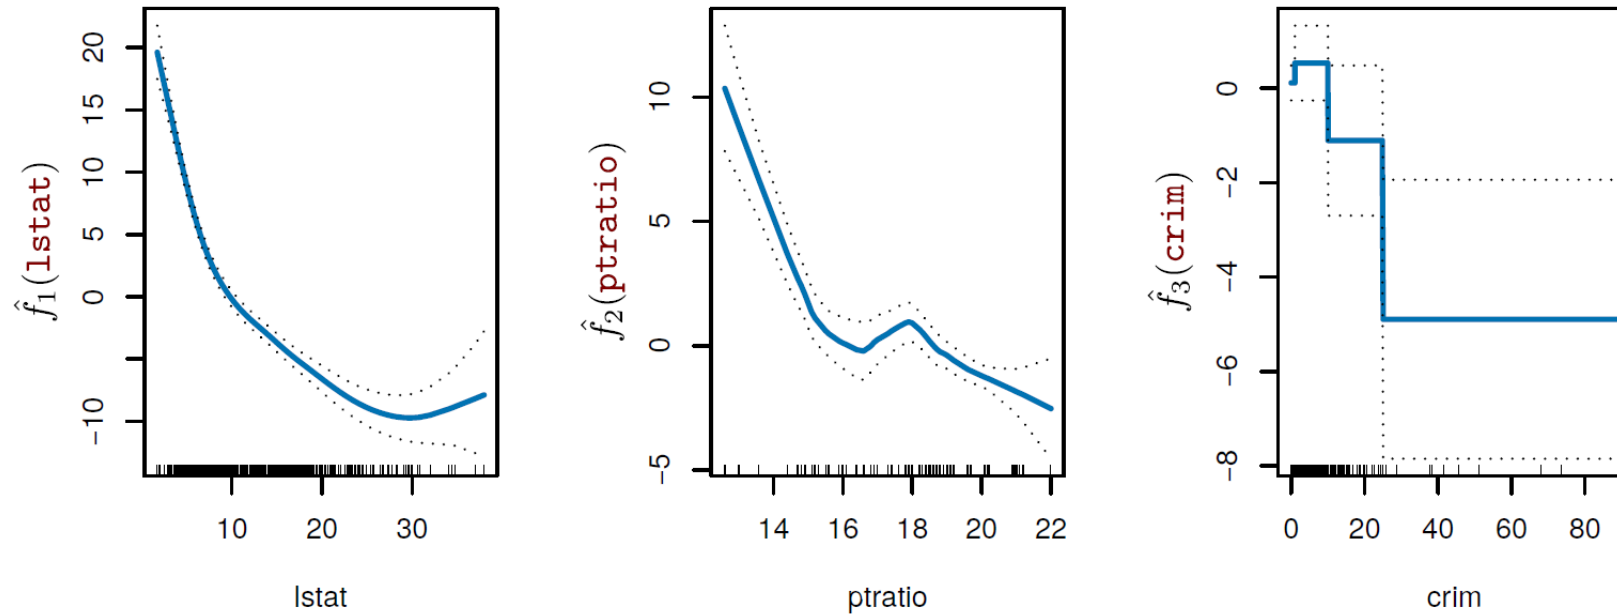
$$Y = \beta_0 + \sum_{j=1}^p \beta_j X_j + \epsilon$$

- We can now extend this to the far more **flexible Additive Model**

$$Y = \beta_0 + \sum_{j=1}^p f_j(X_j) + \epsilon$$

- Each f_j can be **any of the different methods** we just talked about:
Linear term ($\beta_j X_j$), Polynomial, Step Function, Piecewise Polynomial, Degree- k spline, Natural cubic spline, Smoothing spline, Local linear regression fit, ...
- You can mix-and-match different kinds of terms
- The **gam** and **mgcv** packages enable Additive Models in **R**

GAMs: Boston housing data




$$\text{medv} = f_1(\text{lstat}) + f_2(\text{ptratio}) + f_3(\text{crim}) + \epsilon$$

- $f_1(\text{lstat})$ smoothing spline with 5 df
- $f_2(\text{ptratio})$ local linear regression
- $f_3(\text{crim})$ step function with breaks at **crim** = 1, 10, 25

```
gam(medv ~ s(lstat, 5) + lo(ptratio) +  
      cut(crim, breaks = c(-Inf, 1, 10, 25, Inf)),  
     data = Boston)
```


GAMs for classification

$$\log \left(\frac{p(X)}{1 - p(X)} \right) = \beta_0 + f_1(X_1) + f_2(X_2) + \cdots + f_p(X_p).$$



```
gam(I(wage > 250) ~ year + s(age, df = 5) + education, family = binomial)
```