

# **BIOS 635: Principal Components Regression and Partial Least Squares**

Kevin Donovan

3/4/2021

# Review

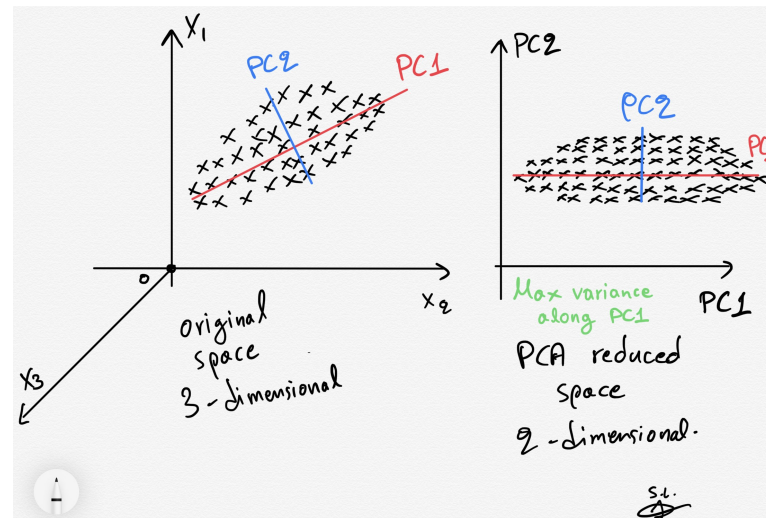
- Midterm assigned on Github Classroom, due on 3/19 by 11 PM EST
- Article evaluation 3 assigned on Github Classroom, due on 3/23 by 11 PM EST
- Last lecture: shrinkage and penalized regression

# Model selection

- **Goal:** Choose/build model parameters and structure to create *optimal* model
- General methods:
  - I. Subset Selection
  2. Shrinkage
  3. **Dimension Reduction**

# Dimension Reduction

- For data with large  $p$ , may want to **reduce predictor set**
  - Reduces chance of overfitting, estimation variance
  - Predictors may be highly correlated, but still important to assess
  - Traditional regression with  $p > n$  not computationally possible
- Idea: **project predictor space into reduced dimensional space**
  - Use these new variables as regression model predictors



# Principal Components Regression

- Original predictors:  $X_1, \dots, X_p$
- New set:  $Z_1, \dots, Z_M$  where  $M < p$ 
  - Suppose  $Z_i$  are **linear combinations** of  $X_1, \dots, X_p$

$$Z_m = \sum_{j=1}^p \phi_{jm} X_j$$

- Regression model:

$$Y_i = \theta_0 + \sum_{m=1}^M \theta_m Z_{im} + \epsilon_i$$

- Dimension of model reduced from  $p + 1$  to  $M + 1$
- **How do we decide on reduced set?**
  1. How big should  $M$  be?
  2. How to estimate  $\{\phi_{jm}\}$ ?

# Principal Components Analysis (PCA)

- PCA creates reduced set of predictors **equal to linear combos** of original set
- Method:
  - *First principal component (PC):*

$$Z_1 = \phi_{11}X_1 + \phi_{21}X_2 + \dots + \phi_{p1}X_p$$

$$\hat{\phi}_1 = \operatorname{argmax}_{\|\phi_1\|=1} [\operatorname{Var}(\phi_{11}X_1 + \phi_{21}X_2 + \dots + \phi_{p1}X_p)]$$

$$= \operatorname{argmax}_{\|\phi_1\|=1} (\phi_{11}^2 + \phi_{21}^2 + \dots + \phi_{p1}^2 + 2\phi_{11}\phi_{21}\rho_{12} + \dots + 2\phi_{11}\phi_{p1}\rho_{1p} + \dots + 2\phi_{p-1,1}\phi_{p1}\rho_{p-1,p})$$

- $\rho_{ij} = \operatorname{Cor}(X_i, X_j)$
- First PC = linear combo containing **maximum amount of variability between predictors**
- Weights  $\phi_{ij}$  call *loadings*

# PCA details


- $\|\phi_1\| = 1 \rightarrow \sum_{j=1}^p \phi_{j1}^2 = 1$
- Above assumes  $X$  predictor matrix is *centered* ( $X_j$  have mean 0) and *scaled* (have variance 1)
- Can then compute PC values for each observation:

$$Z_{i1} = \phi_{11}X_{i1} + \phi_{21}X_{i2} + \dots + \phi_{ip}X_{ip} \text{ for } 1 \leq i \leq n$$

- For next PC  $Z_2$ , same process is done but  $\text{Cor}(Z_1, Z_2) = 0$  assumed
  - $\rightarrow \phi_2$  and  $\phi_1$  are *orthogonal*

# PCA recap

- PCA uses an **orthogonal transformation** to convert predictor set into new set of **uncorrelated** predictors equal to **linear combinations** of the original set
- Can compute new variables for each observation using **loadings**



First principal component:  $\mathbf{Z}_1 = \phi_{11}\mathbf{X}_1 + \phi_{21}\mathbf{X}_2 + \dots + \phi_{p1}\mathbf{X}_p$

PC loading vector:  $\phi_1 = \{\phi_{11}, \phi_{21}, \dots, \phi_{p1}\}^T$ ,  $\sum_{j=1}^p \phi_{j1}^2 = 1$

$\hat{\phi}_1 = \underset{\|\phi_1\|=1}{\operatorname{argmax}} \{ \phi_1^T \mathbf{X}^T \mathbf{X} \phi_1 \}$  maximizes the variance of  $\mathbf{Z}_1$ .

$\phi_1$  is the eigenvector corresponding to the largest eigenvalue of  $\mathbf{X}^T \mathbf{X}$ .

$\mathbf{Z}_2$  is restrained to be uncorrelated with  $\mathbf{Z}_1$ . Equivalently,  $\phi_2$  is orthogonal to  $\phi_1$ .



# PCA example

- Suppose we have a large amount of regional brain activity measures on group of infants
- Want to see if this regional brain activity is related to diagnosis of Autism (ASD)

```
brain_data <- read_csv("../data/IBIS_brain_data_ex.csv")


brain_data <- brain_data %>%
  select("CandID", "RiskGroup", names(brain_data)[grep1("V12", names(brain_data))]) %>%
  select(CandID:Uncinate_R_V12) %>%
  drop_na()

dim(brain_data)
```

```
## [1] 278 120
```

```
# Run PCA on brain variables
pca_brain <- prcomp(brain_data %>%
  select(LeftAmygdala_V12, RightAmygdala_V12,
    PreCG_L_V12:Uncinate_R_V12) %>% as.data.frame(),
  cor=TRUE, scores=TRUE)


# Scree plot
fviz_eig(pca_brain)
```



```
# Plot PC scores
# First will center and scale PC2
data_for_plot <- scale(pca_brain$x)

# Add in diagnosis back
data_for_plot <- data.frame(data_for_plot,
                           "RiskGroup"=brain_data$RiskGroup,
                           "CandID"=brain_data$CandID)


ggplot(data = data.frame(data_for_plot),
       mapping = aes(x=PC1, y=PC2, color=RiskGroup))+
  geom_point()+
  geom_hline(yintercept=0, linetype="dashed")+
  geom_vline(xintercept=0, linetype="dashed")+
  theme_bw()
```



# PCA example

- Explain overall track and field athlete performance using results in many events

```
fviz_pca_var(res.pca,  
  col.var = "contrib", # Color by contributions to the PC  
  gradient.cols = c("#00AFBB", "#E7B800", "#FC4E07"),  
  repel = TRUE # Avoid text overlapping  
)
```



# Sparse PCA

- Some predictors may not be very useful in principal components
  - However, loadings **won't be exactly 0**
  - **Idea:** apply penalized regression idea to PCA
  - Called sparse PCA

First principal component:  $\mathbf{Z}_1 = \phi_{11}\mathbf{X}_1 + \phi_{21}\mathbf{X}_2 + \dots + \phi_{p1}\mathbf{X}_p$

PC loading vector:  $\phi_1 = \{\phi_{11}, \phi_{21}, \dots, \phi_{p1}\}^T$ ,  $\sum_{j=1}^p \phi_{j1}^2 = 1$

$\hat{\phi}_1 = \underset{\|\phi_1\|=1, \|\phi_1\|_0 \leq k}{\operatorname{argmax}} \{ \phi_1^T \mathbf{X}^T \mathbf{X} \phi_1 \}$  maximizes the variance of  $\mathbf{Z}_1$ .

If  $k = p$ , this reduces to ordinary PCA.

# Sparse PCA in R

```
# Run sparse PCA on brain variables
spca_data <- brain_data %>%
  select(LeftAmygdala_V12, RightAmygdala_V12,
         PreCG_L_V12:Uncinate_R_V12) %>% as.data.frame()

sparsepca_brain <- spca(x=spca_data,
                       K=4,
                       para = rep(1, dim(spca_data)[2]),
                       sparse="penalty",
                       use.corr=TRUE)

# Look at Loadings
apply(X=sparsepca_brain$loadings, MARGIN=2, function(x){sum(x==0)})
```

```
## PC1 PC2 PC3 PC4
##   1   2  10   6
```

# Principal Components Regression

- Regression Model:  $Y$ =outcome,  $Z_1, \dots, Z_M$ =predictors
  - $Y$  not used in creation of predictors, thus *unsupervised*
  - $\rightarrow$  PCs may best explain predictors but may **not** be best set at also **predicting response**
- Solution: *partial least squares* (PLS)

# Partial Least Squares

- **Goal:** Find set of  $Z_1, \dots, Z_M$  that best summarizes  $X_1, \dots, X_p$  **and** their relationship to outcome  $Y$
- **Method:**
  1. *Standardize  $p$  predictors*
  2. *For first component  $Z_1$  :*
    - Set each  $\phi_{1j}$  from PCA equal to  $\beta$  from regression of  $Y$  onto  $X_j$
    - Results in  $Z_1 = \sum_{j=1}^p \phi_{1j} X_j$
    - Then compute residuals of regressing  $Y$  onto  $Z_1$
  3. *Repeat for desired number of components*
- **Fit in R:** pls package